

PS3_Babaei

September 30, 2019

```
[1]: # Saharnaz Babaei Balderlou
# Problem Set 3
#-----
# Import required packages
#-----
import pandas as pd # will be used to read .dta file by .read_stata()
import numpy as np
import matplotlib.pyplot as plt # will be used to see the obvious relationship
    ↳ of desired variables in a scatterplot
import matplotlib inline
import math # to use in some equations
from scipy.optimize import minimize # for optimization of Likelihood function
    ↳ (MLE method)
import scipy.optimize as opt
import statsmodels.api as statmod
import scipy.stats as stats
```

```
[2]: #-----
# Read data and prepare to utilize (Part 1 & 2)
#-----
df1 = pd.read_stata('PS3_data.dta')
# My note for later: https://www.shanelynn.ie/using-pandas-dataframe-creating-editing-viewing-data-in-python/
```

```
[3]: print("Dataframe: ")
df1.head()
```

Dataframe:

```
[3]:   id68  year  intid relhh  hannhrs  wannhrs  hlabinc  wlabinc  nochild  \
0     1  1967     1  Head   1200.0   2000.0      NaN     NaN        0
1     2  1967     2  Head     0.0     0.0      NaN     NaN        0
2     3  1967     3  Head     0.0     0.0      NaN     NaN        0
3     4  1967     4  Head   1560.0     0.0      NaN     NaN        6
4     5  1967     5  Head   2500.0   2000.0      NaN     NaN        3

   wrace  ...  redpregovinc  hsex  wsex  age  wage  hpersno  wpersno  hyrsed  \
0    NaN  ...           5614.0   1.0   2.0  52.0  46.0      1.0      2.0    8.0
```

1	NaN	...	0.0	1.0	2.0	56.0	57.0	1.0	2.0	3.0
2	NaN	...	0.0	1.0	2.0	77.0	64.0	1.0	2.0	NaN
3	1.0	...	3280.0	1.0	2.0	45.0	44.0	1.0	2.0	8.0
4	1.0	...	7900.0	1.0	2.0	24.0	22.0	1.0	2.0	10.0

	wyrsed	pce
0	8.0	0.0
1	3.0	0.0
2	3.0	0.0
3	5.0	0.0
4	9.0	0.0

[5 rows x 52 columns]

```
[4]: '''
      hlabinc = annual labor income of the head
      hannhrs = annual hours of the head
      hsex = gender of the head (1 = Male, 2 = Female)
      hrace = race of the head (1 = white, 2 = Black, 3 = Native American, 4 = Asian/
      ↪Pacific Islander, 5 = Hispanic, 6,7 = Other)
      age = age of the head
      hyrsed = years of education of the head
      '''
      print("Data Statistics:")
      df1.describe()
```

Data Statistics:

```
[4]:
```

	id68	year	intid	hannhrs \
count	123786.000000	123786.000000	123786.000000	123786.000000
mean	1494.639475	1984.831273	3271.379429	1679.269897
std	838.901790	9.836212	2277.056058	1061.704712
min	1.000000	1967.000000	1.000000	0.000000
25%	772.000000	1977.000000	1444.000000	832.000000
50%	1517.000000	1985.000000	2984.000000	1976.000000
75%	2224.000000	1993.000000	4763.000000	2350.000000
max	2930.000000	2002.000000	16968.000000	7800.000000

	wannhrs	hlabinc	wlabinc	nochild \
count	123786.000000	9.023300e+04	48496.000000	123786.000000
mean	633.026917	4.211505e+04	22026.289062	0.843771
std	878.422791	4.670424e+04	21336.107422	1.182829
min	0.000000	6.353981e-01	1.192780	0.000000
25%	0.000000	1.979858e+04	8016.247070	0.000000
50%	0.000000	3.460022e+04	18122.412109	0.000000
75%	1454.000000	5.267309e+04	30256.060547	2.000000
max	5840.000000	3.771521e+06	856942.062500	11.000000

	wrace	hrace	...	redpregovinc	hsex	wsex	\
count	90603.000000	123656.000000	...	1.237860e+05	123786.000000	80758.0	
mean	1.098220	1.129731	...	3.012258e+04	1.233072	2.0	
std	0.356161	0.394627	...	4.588795e+04	0.422940	0.0	
min	1.000000	1.000000	...	-1.324040e+05	1.000000	2.0	
25%	1.000000	1.000000	...	7.700000e+03	1.000000	2.0	
50%	1.000000	1.000000	...	1.900000e+04	1.000000	2.0	
75%	1.000000	1.000000	...	3.910775e+04	1.000000	2.0	
max	8.000000	8.000000	...	3.660000e+06	2.000000	2.0	


	age	wage	hpersno	wpersno	\
count	123786.000000	80758.000000	123786.000000	80758.000000	
mean	45.545547	41.390785	39.620201	55.346169	
std	17.623671	14.786721	69.003265	77.864296	
min	16.000000	13.000000	1.000000	1.000000	
25%	31.000000	29.000000	1.000000	2.000000	
50%	42.000000	39.000000	3.000000	3.000000	
75%	58.000000	51.000000	22.000000	170.000000	
max	102.000000	95.000000	227.000000	231.000000	

	hyrsed	wyrsed	pce
count	122809.000000	80091.000000	123786.000000
mean	12.666091	12.720081	0.557690
std	2.917721	2.422607	0.265198
min	1.000000	1.000000	0.000000
25%	12.000000	12.000000	0.362158
50%	12.000000	12.000000	0.599887
75%	15.000000	14.000000	0.786908
max	17.000000	17.000000	0.928007

[8 rows x 51 columns]

```
[5]: print("Scatterplot between annual labor income of the head and years of
      ↪education of the head")
plt.style.use('seaborn')
df1.plot(x = 'hyrsed', y = 'hlabinc', kind = 'scatter')
plt.show()
plt.savefig('scat_inc_edu.png')
```

Scatterplot between annual labor income of the head and years of education of the head



<Figure size 576x396 with 0 Axes>

```
[6]: #drop missing values
df1_subset = df1.dropna(how = 'any', subset = ['hlabinc', 'hannhrs', 'hsex', 'h
→ 'hrace', 'age', 'hyrsed'])
df1_subset.describe()
```

```
[6]:
```

	id68	year	intid	hannhrs	wannhrs	\
count	89688.000000	89688.000000	89688.000000	89688.000000	89688.000000	
mean	1510.782992	1986.315338	3515.175007	2067.510254	763.457153	
std	834.667982	8.791094	2314.034043	756.294312	913.115479	
min	1.000000	1971.000000	1.000000	0.000000	0.000000	
25%	783.000000	1979.000000	1673.000000	1800.000000	0.000000	
50%	1542.000000	1986.000000	3321.000000	2064.000000	74.000000	
75%	2240.000000	1993.000000	5058.000000	2453.000000	1700.000000	
max	2930.000000	2002.000000	16968.000000	7800.000000	5840.000000	

	hlabinc	wlabinc	nochild	wrace	hrace	\
count	8.968800e+04	45338.000000	89688.000000	73835.000000	89688.000000	
mean	4.211382e+04	21914.417969	0.949737	1.096932	1.123695	
std	4.675834e+04	20676.205078	1.168268	0.354777	0.390376	
min	6.353981e-01	1.192780	0.000000	1.000000	1.000000	
25%	1.976367e+04	8042.066406	0.000000	1.000000	1.000000	

50%	3.460022e+04	18120.386719	0.000000	1.000000	1.000000
75%	5.267309e+04	30172.169922	2.000000	1.000000	1.000000
max	3.771521e+06	685266.750000	11.000000	8.000000	3.000000

	...	redpregovinc	hsex	wsex	age	wage \
count	...	8.968800e+04	89688.000000	62705.0	89688.000000	62705.000000
mean	...	3.774112e+04	1.179935	2.0	40.012321	38.289116
std	...	4.853038e+04	0.384065	0.0	13.287443	12.255554
min	...	-9.359900e+04	1.000000	2.0	17.000000	14.000000
25%	...	1.400000e+04	1.000000	2.0	29.000000	29.000000
50%	...	2.640000e+04	1.000000	2.0	38.000000	36.000000
75%	...	4.744325e+04	1.000000	2.0	49.000000	46.000000
max	...	3.660000e+06	2.000000	2.0	95.000000	93.000000

	hpersno	wpersno	hyrsed	wyrsed	pce
count	89688.000000	62705.000000	89688.000000	62244.000000	89688.000000
mean	48.742619	65.553192	13.228726	13.029015	0.609756
std	73.709778	81.090965	2.526992	2.225578	0.208654
min	1.000000	1.000000	1.000000	1.000000	0.247121
25%	1.000000	2.000000	12.000000	12.000000	0.421747
50%	4.000000	4.000000	12.000000	12.000000	0.614522
75%	170.000000	170.000000	16.000000	15.000000	0.786908
max	227.000000	231.000000	17.000000	17.000000	0.928007

[8 rows x 51 columns]

```
[7]: # select male heads of HH whose age is between 25 & 60 (included!), and wage >= $7/hr
      →$7/hr
df2 = df1_subset[['id68', 'year', 'intid', 'hlabinc', 'hannhrs', 'hsex', 'hrace', 'age', 'hyrsed']]
df2['annhrs'] = df2['hannhrs'].where(df2['hannhrs']>0)
df2['hrwage'] = df2['hlabinc']/df2['annhrs'] # compute hourly wage
df2 = df2[(df2.hsex == 1.0) & (df2.age >= 25) & (df2.age <= 60) & (df2.hrwage >= 7)] #Part 1
df2['ln_hr wage'] = np.log(df2['hrwage']) #log of wages
df2.describe()
```

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:3:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

This is separate from the ipykernel package so we can avoid doing imports until

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:4:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>
after removing the cwd from sys.path.

```
[7]:
```

	id68	year	intid	hlabinc	hannhrs \
count	57062.000000	57062.000000	57062.000000	5.706200e+04	57062.000000
mean	1507.470558	1986.575672	3480.375311	5.282805e+04	2228.557617
std	828.361439	8.712311	2253.229068	5.236477e+04	620.054077
min	1.000000	1971.000000	1.000000	1.666980e+01	2.000000
25%	782.000000	1979.000000	1690.000000	3.037345e+04	1952.000000
50%	1542.000000	1987.000000	3296.000000	4.382858e+04	2160.000000
75%	2225.000000	1994.000000	5002.000000	6.138495e+04	2519.000000
max	2930.000000	2002.000000	16968.000000	3.771521e+06	5840.000000

	hsex	hrace	age	hyrsed	annhrs \
count	57062.0	57062.000000	57062.000000	57062.000000	57062.000000
mean	1.0	1.101416	39.243629	13.529967	2228.557617
std	0.0	0.369015	9.578858	2.450013	620.054077
min	1.0	1.000000	25.000000	1.000000	2.000000
25%	1.0	1.000000	31.000000	12.000000	1952.000000
50%	1.0	1.000000	38.000000	13.000000	2160.000000
75%	1.0	1.000000	47.000000	16.000000	2519.000000
max	1.0	3.000000	60.000000	17.000000	5840.000000

	hrwage	ln_hrwage
count	57062.000000	57062.000000
mean	24.320921	3.010804
std	25.209909	0.544096
min	7.000252	1.945946
25%	13.950662	2.635527
50%	19.914677	2.991457
75%	27.790929	3.324710
max	1717.330322	7.448526

```
[8]: # This is a note for me: https://stackoverflow.com/questions/11587782/
      ↪creating-dummy-variables-in-pandas-for-python
race_dummy = pd.get_dummies(df2['hrace']) # Defining race dummies (There is no
      ↪Hispanic individual in dataset)
df2['constant'] = 1
data = pd.concat([df2, race_dummy], axis = 1)
data.rename(columns = {1.0: 'White', 2.0: 'Black', 3.0: 'Others'}, inplace =
      ↪True) #final data "data"; ready to estimate model
data.head()
```

```
[8]:
```

	id68	year	intid	hlabinc	hannhrs	hsex	hrace	age	hyrsed	\
11161	402	1971	1	62928.707031	1523.0	1.0	1.0	51.0	12.0	
11164	461	1971	4	22660.970703	2010.0	1.0	1.0	55.0	5.0	
11166	1126	1971	8	29337.865234	2860.0	1.0	1.0	25.0	16.0	
11173	284	1971	20	76885.437500	2400.0	1.0	1.0	39.0	16.0	
11175	50	1971	29	31968.156250	3164.0	1.0	1.0	36.0	12.0	

	annhrs	hrwage	ln_hrwage	constant	White	Black	Others
11161	1523.0	41.318916	3.721320	1	1	0	0
11164	2010.0	11.274115	2.422509	1	1	0	0
11166	2860.0	10.257995	2.328057	1	1	0	0
11173	2400.0	32.035599	3.466848	1	1	0	0
11175	3164.0	10.103716	2.312903	1	1	0	0

```
[9]: data.dtypes
```

```
[9]: id68          int16
year           int16
intid          int16
hlabinc        float32
hannhrs        float32
hsex           float32
hrace          float64
age            float32
hyrsed         float32
annhrs         float32
hrwage         float32
ln_hrwage      float32
constant       int64
White          uint8
Black          uint8
Others         uint8
dtype: object
```

```
[10]: np.save("data", data)

# Separate years for estimation
data1971 = data[data['year'] == 1971]
np.save("data1971", data1971)
data1980 = data[data['year'] == 1980]
np.save("data1980", data1980)
data1990 = data[data['year'] == 1990]
np.save("data1990", data1990)
data2000 = data[data['year'] == 2000]
np.save("data2000", data2000)
```

```
[11]: #-----
# ML estimation (Part 3)
#-----
```

```

'''
ln(w_it) = alpha + beta_1 * Educ_it + beta_2 * Age_it + beta_3 * Black_it +
↳ beta_4 * Hispanic_it + beta_5 * OtherRace_it + epsilon_it
w_it = wage of individual i in survey year t
Educ_it = education in years
Age_it = age in years
Black_it, Hispanic_it, OtherRace_it = dummy variables for race = Black,
↳ Hispanic, Not (belongs to {White, Black, Hispanic})
'''

# Define my objective function
def myLL(params, t):
    # Coeff.s
    beta0, beta1, beta2, beta3, beta4, sigma = params
    beta = np.array([beta0, beta1, beta2, beta3, beta4])
    n = len(t)
    # Independent variables matrix (No Hispanic)
    x0 = np.array(t['constant']).astype('float')
    x1 = np.array(t['hyrsed']).astype('float')
    x2 = np.array(t['age']).astype('float')
    x3 = np.array(t['Black']).astype('float')
    x4 = np.array(t['Others']).astype('float')
    X = np.column_stack((x0, x1, x2, x3, x4))
    # Dependent variable matrix
    y = np.array(t['ln_hrwage']).astype('float')
    y_bar = np.dot(X, beta)
    ll = -(n/2)*np.log(2*np.pi) - (n/2)*np.log(sigma**2) - (1/
↳ (2*sigma**2))*((y-y_bar).T @ (y - y_bar)))
    return (-ll)

```

```

[12]: # MLE; 'Nelder-Mead'
nbeta = 5
beta = np.zeros(nbeta)
beta0 = 0.1
beta1 = 0.1
beta2 = 0.1
beta3 = 0.1
beta4 = 0.1
sigma = 0.1
beta = [beta0, beta1, beta2, beta3, beta4]
params = [beta0, beta1, beta2, beta3, beta4, sigma]

bounds = ((1e-10, None), (None, None), (None, None), (None, None), (None, None),
↳ (None, None))

res_NM = opt.minimize(myLL, params, args=(data), method='Nelder-Mead',
↳ bounds=bounds)

```



```

print("MLE coefficients: ", "Total dataset")
print("=====")
print(res_NM)
print("_____")
res71_NM = opt.minimize(myLL, params, args=(data1971), method='Nelder-Mead',
    ↳ bounds=bounds)
print("MLE coefficients: ", "year == 1971")
print("=====")
print(res71_NM)
print("_____")
res80_NM = opt.minimize(myLL, params, args=(data1980), method='Nelder-Mead',
    ↳ bounds=bounds)
print("MLE coefficients: ", "year == 1980")
print("=====")
print(res80_NM)
print("_____")
res90_NM = opt.minimize(myLL, params, args=(data1990), method='Nelder-Mead',
    ↳ bounds=bounds)
print("MLE coefficients: ", "year == 1990")
print("=====")
print(res90_NM)
print("_____")
res2000_NM = opt.minimize(myLL, params, args=(data2000), method='Nelder-Mead',
    ↳ bounds=bounds)
print("MLE coefficients: ", "year == 2000")
print("=====")
print(res2000_NM)
print("_____")
'''
years = ["data1971.npy", "data1980.npy", "data1990.npy", "data2000.npy", "data.
    ↳ .npy"]
for t in years:
    res_MLE = opt.minimize(myLL, params, args=(t), method='Nelder-Mead',
    ↳ bounds=bounds)
    print("MLE coefficients: ", t)
    print(res_MLE)
    print("_____")
#### I tried this loop so many times and attempted to troubleshoot but finally
    ↳ I got this error: "string indices must be integers"
'''

```

C:\ProgramData\Anaconda3\lib\site-packages\scipy\optimize_minimize.py:517:
RuntimeWarning: Method Nelder-Mead cannot handle constraints nor bounds.
RuntimeWarning)

MLE coefficients: Total dataset
=====

```

final_simplex: (array([[ 1.39309627,  0.0781139 ,  0.0145177 , -0.16239436,
0.01108259,
    0.48968236],
    [ 1.3930527 ,  0.0781193 ,  0.01451702, -0.1624352 ,  0.01106685,
    0.4896769 ],
    [ 1.39309914,  0.07811356,  0.01451801, -0.16244586,  0.0109965 ,
    0.48968791],
    [ 1.39307198,  0.07811841,  0.01451689, -0.16244504,  0.01111231,
    0.48968793],
    [ 1.39313048,  0.07811415,  0.01451716, -0.16247289,  0.01103884,
    0.48967662],
    [ 1.3930366 ,  0.07812007,  0.01451746, -0.16232213,  0.01103599,
    0.48968893],
    [ 1.39317897,  0.07811236,  0.01451641, -0.16240374,  0.01101757,
    0.48968747]]), array([40225.31420626, 40225.31421319, 40225.31422174,
40225.31422664,
    40225.31424347, 40225.3142488 , 40225.3142569 ]))
    fun: 40225.31420626389
    message: 'Optimization terminated successfully.'
    nfev: 1162
    nit: 749
    status: 0
    success: True
    x: array([ 1.39309627,  0.0781139 ,  0.0145177 , -0.16239436,
0.01108259,
    0.48968236])

```

```

-----
MLE coefficients:  year == 1971
=====

```

```

final_simplex: (array([[ 1.47367052,  0.06973573,  0.01578288, -0.2340242 ,
-0.42911197,
    0.41157107],
    [ 1.47361046,  0.06974115,  0.01578231, -0.23397205, -0.42911159,
    0.4115673 ],
    [ 1.47358856,  0.06973973,  0.01578298, -0.23403964, -0.42908831,
    0.41157356],
    [ 1.4736112 ,  0.06973677,  0.01578324, -0.23399855, -0.42909645,
    0.41156275],
    [ 1.47371123,  0.06973561,  0.01578128, -0.23405239, -0.42911123,
    0.41156895],
    [ 1.47367534,  0.06973616,  0.01578253, -0.23405703, -0.42909264,
    0.41156107],
    [ 1.47371518,  0.06973268,  0.01578194, -0.23400247, -0.42914202,
    0.41158344]]), array([752.78870883, 752.78870947, 752.78871087,
752.78871124,
    752.78871183, 752.7887139 , 752.78871442]))
    fun: 752.7887088280063
    message: 'Optimization terminated successfully.'

```

```

nfev: 786
nit: 496
status: 0
success: True
x: array([ 1.47367052,  0.06973573,  0.01578288, -0.2340242 ,
-0.42911197,
0.41157107])

-----
MLE coefficients:  year == 1980
=====
final_simplex: (array([[ 1.6132899 ,  0.0675381 ,  0.01269899, -0.10263128,
0.0136228 ,
0.44926559],
[ 1.61305248,  0.06756315,  0.01269762, -0.10267639,  0.01366701,
0.44928481],
[ 1.61340895,  0.06753724,  0.01269743, -0.10258163,  0.01316331,
0.4492538 ],
[ 1.61323268,  0.06753902,  0.01269958, -0.10288796,  0.01379478,
0.44921281],
[ 1.61374566,  0.06752382,  0.01269221, -0.10279164,  0.01359016,
0.44925456],
[ 1.61356332,  0.06752716,  0.01269722, -0.10289006,  0.01380762,
0.44925926],
[ 1.61340448,  0.06753553,  0.01269851, -0.10301816,  0.01315615,
0.4492721 ]]), array([1148.39323838, 1148.39324439, 1148.39325561,
1148.39325705,
1148.39326198, 1148.39326495, 1148.39326782])))
fun: 1148.39323837555
message: 'Maximum number of function evaluations has been exceeded.'
nfev: 1201
nit: 774
status: 1
success: False
x: array([ 1.6132899 ,  0.0675381 ,  0.01269899, -0.10263128,
0.0136228 ,
0.44926559])

-----
MLE coefficients:  year == 1990
=====
final_simplex: (array([[ 0.66841963,  0.11863582,  0.01749401, -0.14739724,
-0.65477016,
0.4893061 ],
[ 0.66835602,  0.11864754,  0.0174916 , -0.14732871, -0.65471547,
0.4893005 ],
[ 0.66851798,  0.1186424 ,  0.01748935, -0.14746815, -0.65485297,
0.48929941],
[ 0.66837051,  0.11864686,  0.01749027, -0.14739821, -0.65472676,
0.48930142],

```

```

[ 0.66838604, 0.11864412, 0.01749209, -0.14745596, -0.65473298,
 0.48929077],
[ 0.66840779, 0.11864285, 0.01749167, -0.14741798, -0.65476171,
 0.4893055 ],
[ 0.66846521, 0.11864223, 0.01749013, -0.14735539, -0.65480545,
 0.48929722]]), array([1430.23330055, 1430.23330579, 1430.23330615,
1430.23330668,
1430.23330737, 1430.23330866, 1430.23331002]))
    fun: 1430.233300554175
    message: 'Optimization terminated successfully.'
    nfev: 852
    nit: 547
    status: 0
    success: True
    x: array([ 0.66841963, 0.11863582, 0.01749401, -0.14739724,
-0.65477016,
    0.4893061 ]))

-----
MLE coefficients: year == 2000
=====
    final_simplex: (array([[ 1.16169112, 0.10915564, 0.01099325, -0.24608949,
-0.06074079,
    0.5395586 ],
[ 1.16172663, 0.10915285, 0.01099324, -0.24602461, -0.06069615,
    0.53956308],
[ 1.16176407, 0.10914997, 0.01099332, -0.2460459 , -0.06071587,
    0.53956483],
[ 1.16165857, 0.10915336, 0.01099491, -0.24603163, -0.06075654,
    0.53955864],
[ 1.16175584, 0.10915004, 0.01099372, -0.24609096, -0.06069359,
    0.53956002],
[ 1.16161566, 0.10915489, 0.01099514, -0.24607217, -0.06067408,
    0.53955494],
[ 1.1616349 , 0.10915557, 0.01099461, -0.24605415, -0.06071328,
    0.53954611]]), array([2068.985145 , 2068.98514519, 2068.98514543,
2068.98514579,
2068.98514608, 2068.98514614, 2068.98514625]))
    fun: 2068.985144996586
    message: 'Optimization terminated successfully.'
    nfev: 1143
    nit: 742
    status: 0
    success: True
    x: array([ 1.16169112, 0.10915564, 0.01099325, -0.24608949,
-0.06074079,
    0.5395586 ]))

-----

```

```
[12]: '\nyears = ["data1971.npy", "data1980.npy", "data1990.npy", "data2000.npy",
"data.npy"]\nfor t in years:\n    res_MLE = opt.minimize(myLL, params, args=(t),
method='\Nelder-Mead', bounds=bounds)\n    print("MLE coefficients: ", t)\n
print(res_MLE)\n    print("_____")\n#### I tried
this loop so many times and attempted to troubleshoot but finally I got this
error: "string indices must be integers"\n'
```

```
[13]: # MLE; 'L-BFGS-B'
nbeta = 5
beta = np.zeros(nbeta)
beta0 = 0.1
beta1 = 0.1
beta2 = 0.1
beta3 = 0.1
beta4 = 0.1
sigma = 0.1
beta = [beta0, beta1, beta2, beta3, beta4]
params = [beta0, beta1, beta2, beta3, beta4, sigma]

bounds = ((1e-10, None), (None, None), (None, None), (None, None), (None, None),
→(None, None))

res_B = opt.minimize(myLL, params, args=(data), method='L-BFGS-B',
→bounds=bounds)
print("MLE coefficients: ", "Total dataset")
print("=====")
print(res_B)
print("_____")
res71_B = opt.minimize(myLL, params, args=(data1971), method='L-BFGS-B',
→bounds=bounds)
print("MLE coefficients: ", "year == 1971")
print("=====")
print(res71_B)
print("_____")
res80_B = opt.minimize(myLL, params, args=(data1980), method='L-BFGS-B',
→bounds=bounds)
print("MLE coefficients: ", "year == 1980")
print("=====")
print(res80_B)
print("_____")
res90_B = opt.minimize(myLL, params, args=(data1990), method='L-BFGS-B',
→bounds=bounds)
print("MLE coefficients: ", "year == 1990")
print("=====")
print(res90_B)
print("_____")
```

```

res2000_B = opt.minimize(myLL, params, args=(data2000), method='L-BFGS-B',
    ↪ bounds=bounds)
print("MLE coefficients: ", "year == 2000")
print("=====")
print(res2000_B)
print("-----")

```

MLE coefficients: Total dataset

```

=====
      fun: 40225.31422021307
    hess_inv: <6x6 LbfgsInvHessProduct with dtype=float64>
       jac: array([ 1.09284883, 11.08637662, 13.14037945,  0.50786184,
-0.14697434,
      2.48692231])
    message: b'CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH'
       nfev: 469
        nit: 57
     status: 0
    success: True
         x: array([ 1.39310199,  0.07811655,  0.01451693, -0.1623717 ,  0.0110262
,
      0.48968773])

```

MLE coefficients: year == 1971

```

=====
      fun: 728.0628698968592
    hess_inv: <6x6 LbfgsInvHessProduct with dtype=float64>
       jac: array([-5.36374500e-02, -4.52814675e-01, -2.12465920e+00,
3.09682946e-02,
      -7.27595761e-04,  9.65997060e-02])
    message: b'CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH'
       nfev: 602
        nit: 76
     status: 0
    success: True
         x: array([ 1.55089935,  0.0668825 ,  0.01439164, -0.16381197,  0.0306993
,
      0.41010486])

```

MLE coefficients: year == 1980

```

=====
      fun: 1148.3932188882504
    hess_inv: <6x6 LbfgsInvHessProduct with dtype=float64>
       jac: array([-0.00618456, -0.013506 , -0.2509978 ,  0.01757599,
-0.01077751,
      -0.02992238])
    message: b'CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH'

```

```

    nfev: 448
    nit: 52
    status: 0
    success: True
    x: array([ 1.61306755,  0.06755729,  0.01269855, -0.10270202,
0.01346053,
    0.44924099])

```

```

-----
MLE coefficients:  year == 1990
=====

```

```

    fun: 1393.8821505146818
    hess_inv: <6x6 LbfgsInvHessProduct with dtype=float64>
    jac: array([-1.40516931e-02, -2.04931894e-01, -5.94741323e-01,
5.91171556e-04,
    -2.95585778e-04, -1.26192390e-02])
    message: b'CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH'
    nfev: 511
    nit: 64
    status: 0
    success: True
    x: array([ 1.11859519,  0.0975577 ,  0.01346543, -0.17202183,
-0.05971275,
    0.4835987  ])

```

```

-----
MLE coefficients:  year == 2000
=====

```

```

    fun: 2068.985144417821
    hess_inv: <6x6 LbfgsInvHessProduct with dtype=float64>
    jac: array([ 0.00245564,  0.03237801,  0.03697096, -0.00118234,
0.00172804,
    -0.0125965  ])
    message: b'CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH'
    nfev: 378
    nit: 47
    status: 0
    success: True
    x: array([ 1.16170054,  0.1091542 ,  0.01099335, -0.24604792,
-0.06072719,
    0.53955657])

```

```

[14]: # MLE; 'SLSQP'
nbeta = 5
beta = np.zeros(nbeta)
beta0 = 1.20
beta1 = 0.11
beta2 = 0.01

```

```

beta3 = 0.20
beta4 = 0.01
sigma = 0.50
beta = [beta0, beta1, beta2, beta3, beta4]
params = [beta0, beta1, beta2, beta3, beta4, sigma]

bounds = ((1e-10, None), (None, None), (None, None), (None, None), (None, None),
→(None, None))

res_S = opt.minimize(myLL, params, args=(data), method='SLSQP', bounds=bounds)
print("MLE coefficients: ", "Total dataset")
print("=====")
print(res_S)
print(" _____")
res71_S = opt.minimize(myLL, params, args=(data1971), method='SLSQP',
→bounds=bounds)
print("MLE coefficients: ", "year == 1971")
print("=====")
print(res71_S)
print(" _____")
res80_S = opt.minimize(myLL, params, args=(data1980), method='SLSQP',
→bounds=bounds)
print("MLE coefficients: ", "year == 1980")
print("=====")
print(res80_S)
print(" _____")
res90_S = opt.minimize(myLL, params, args=(data1990), method='SLSQP',
→bounds=bounds)
print("MLE coefficients: ", "year == 1990")
print("=====")
print(res90_S)
print(" _____")
res2000_S = opt.minimize(myLL, params, args=(data2000), method='SLSQP',
→bounds=bounds)
print("MLE coefficients: ", "year == 2000")
print("=====")
print(res2000_S)
print(" _____")

```

MLE coefficients: Total dataset

=====

fun: 40225.314173424194

jac: array([0.0078125 , 0.08886719, 0.16503906, -0.00683594, 0.

,

0.11083984])

message: 'Optimization terminated successfully.'

nfev: 140


```

        nit: 12
        njev: 12
        status: 0
        success: True
        x: array([ 1.39301236,  0.07811904,  0.01451826, -0.16240252,
0.01105784,
        0.48968273])

-----
MLE coefficients:  year == 1971
=====
        fun: 728.062867806588
        jac: array([-0.06733704, -0.85600281, -2.66024017,  0.01177979,
-0.00302124,
        0.07466125])
        message: 'Optimization terminated successfully.'
        nfev: 125
        nit: 11
        njev: 11
        status: 0
        success: True
        x: array([ 1.55097459,  0.06687866,  0.01439123, -0.1638947 ,
0.03068816,
        0.41009816])

-----
MLE coefficients:  year == 1980
=====
        fun: 1148.3932185943872
        jac: array([-0.0244751 , -0.30895996, -0.92105103,  0.01419067,
-0.00190735,
        0.05471802])
        message: 'Optimization terminated successfully.'
        nfev: 120
        nit: 11
        njev: 11
        status: 0
        success: True
        x: array([ 1.61308089,  0.06755625,  0.01269849, -0.10270663,
0.01350651,
        0.44924558])

-----
MLE coefficients:  year == 1990
=====
        fun: 1393.8821506055688
        jac: array([-0.02429199, -0.3656311 , -0.88607788, -0.00256348,  0.
,
        0.04708862])
        message: 'Optimization terminated successfully.'
        nfev: 117

```

```

        nit: 11
        njev: 11
        status: 0
        success: True
        x: array([ 1.11859812,  0.0975571 ,  0.01346555, -0.17202805, -0.0597097
,
        0.48360217])

-----
MLE coefficients:  year == 2000
=====
        fun: 2068.9851444238884
        jac: array([ 1.58691406e-02,  2.22412109e-01,  6.81030273e-01,
2.07519531e-03,
        4.27246094e-04, -1.78222656e-02])
        message: 'Optimization terminated successfully.'
        nfev: 115
        nit: 10
        njev: 10
        status: 0
        success: True
        x: array([ 1.16169916,  0.1091542 ,  0.01099342, -0.24604219,
-0.06073273,
        0.53955628])

-----

```

[15]: *# Comparing the MLE coefficients for 3 different methods:*

```

print("All data; Nelder-Mead: ", res_NM.x)
print("1971; Nelder-Mead: ", res71_NM.x)
print("1980; Nelder-Mead: ", res80_NM.x)
print("1990; Nelder-Mead: ", res90_NM.x)
print("2000; Nelder-Mead: ", res2000_NM.x)
print("-----")
print("All data; L-BFGS-B: ", res_B.x)
print("1971; L-BFGS-B: ", res71_B.x)
print("1980; L-BFGS-B: ", res80_B.x)
print("1990; L-BFGS-B: ", res90_B.x)
print("2000; L-BFGS-B: ", res2000_B.x)
print("-----")
print("All data; SLSQP: ", res_S.x)
print("1971; L-BFGS-B: ", res71_S.x)
print("1980; L-BFGS-B: ", res80_S.x)
print("1990; L-BFGS-B: ", res90_S.x)
print("2000; L-BFGS-B: ", res2000_S.x)

```

```

All data; Nelder-Mead: [ 1.39309627  0.0781139   0.0145177  -0.16239436
0.01108259  0.48968236]
1971; Nelder-Mead: [ 1.47367052  0.06973573  0.01578288 -0.2340242  -0.42911197

```

```

0.41157107]
1980; Nelder-Mead: [ 1.6132899  0.0675381  0.01269899 -0.10263128  0.0136228
0.44926559]
1990; Nelder-Mead: [ 0.66841963  0.11863582  0.01749401 -0.14739724 -0.65477016
0.4893061 ]
2000; Nelder-Mead: [ 1.16169112  0.10915564  0.01099325 -0.24608949 -0.06074079
0.5395586 ]

-----
All data; L-BFGS-B: [ 1.39310199  0.07811655  0.01451693 -0.1623717  0.0110262
0.48968773]
1971; L-BFGS-B: [ 1.55089935  0.0668825  0.01439164 -0.16381197  0.0306993
0.41010486]
1980; L-BFGS-B: [ 1.61306755  0.06755729  0.01269855 -0.10270202  0.01346053
0.44924099]
1990; L-BFGS-B: [ 1.11859519  0.0975577  0.01346543 -0.17202183 -0.05971275
0.4835987 ]
2000; L-BFGS-B: [ 1.16170054  0.1091542  0.01099335 -0.24604792 -0.06072719
0.53955657]

-----
All data; SLSQP: [ 1.39301236  0.07811904  0.01451826 -0.16240252  0.01105784
0.48968273]
1971; L-BFGS-B: [ 1.55097459  0.06687866  0.01439123 -0.1638947  0.03068816
0.41009816]
1980; L-BFGS-B: [ 1.61308089  0.06755625  0.01269849 -0.10270663  0.01350651
0.44924558]
1990; L-BFGS-B: [ 1.11859812  0.0975571  0.01346555 -0.17202805 -0.0597097
0.48360217]
2000; L-BFGS-B: [ 1.16169916  0.1091542  0.01099342 -0.24604219 -0.06073273
0.53955628]

```

```

[16]: # OLS

#https://lectures.quantecon.org/py/ols.html
# to estimate each year separately
res_OLS = statmod.OLS(endog=data['ln_hrwage'], exog=data[['constant', 'hyrsed',
→'age', 'Black', 'Others']]).fit()
res71_OLS = statmod.OLS(endog=data1971['ln_hrwage'], exog=data1971[['constant',
→'hyrsed', 'age', 'Black', 'Others']]).fit()
res80_OLS = statmod.OLS(endog=data1980['ln_hrwage'], exog=data1980[['constant',
→'hyrsed', 'age', 'Black', 'Others']]).fit()
res90_OLS = statmod.OLS(endog=data1990['ln_hrwage'], exog=data1990[['constant',
→'hyrsed', 'age', 'Black', 'Others']]).fit()
res2000_OLS = statmod.OLS(endog=data2000['ln_hrwage'],
→exog=data2000[['constant', 'hyrsed', 'age', 'Black', 'Others']]).fit()

print('OLS; Full Sample')
print('=====')

```

```

print(res_OLS.summary())
print('~~~~~')
print('OLS; year == 1971')
print('=====')
print(res71_OLS.summary())
print('~~~~~')
print('OLS; year == 1980')
print('=====')
print(res80_OLS.summary())
print('~~~~~')
print('OLS; year == 1990')
print('=====')
print(res90_OLS.summary())
print('~~~~~')
print('OLS; year == 2000')
print('=====')
print(res2000_OLS.summary())
print('~~~~~')
'''
years = ['data1971.npy', 'data1980.npy', 'data1990.npy', 'data2000.npy', 'data.
→npy']
for t in years:
    print('OLS for', t)
    print('=====')
    print(statmod.OLS(endog=t['ln_hrwage'], exog=t[['constant', 'hyrsed',
→'age', 'Black', 'Others']]).fit().summary())
    □
→print('-----')
'''

```

OLS; Full Sample

=====

OLS Regression Results

```

=====
Dep. Variable:          ln_hrwage    R-squared:               0.190
Model:                  OLS          Adj. R-squared:          0.190
Method:                 Least Squares    F-statistic:             3346.
Date:                  Mon, 30 Sep 2019    Prob (F-statistic):       0.00
Time:                  20:41:36          Log-Likelihood:          -40225.
No. Observations:      57062           AIC:                    8.046e+04
Df Residuals:          57057           BIC:                    8.051e+04
Df Model:               4
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
constant	1.3930	0.015	94.165	0.000	1.364	1.422

hyrsed	0.0781	0.001	92.558	0.000	0.076	0.080
age	0.0145	0.000	67.707	0.000	0.014	0.015
Black	-0.1624	0.009	-18.129	0.000	-0.180	-0.145
Others	0.0111	0.014	0.800	0.424	-0.016	0.038

Omnibus:	5824.662	Durbin-Watson:	1.937
Prob(Omnibus):	0.000	Jarque-Bera (JB):	16025.805
Skew:	0.572	Prob(JB):	0.00
Kurtosis:	5.330	Cond. No.	310.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS; year == 1971

OLS Regression Results

Dep. Variable:	ln_hrwage	R-squared:	0.244
Model:	OLS	Adj. R-squared:	0.241
Method:	Least Squares	F-statistic:	110.7
Date:	Mon, 30 Sep 2019	Prob (F-statistic):	7.42e-82
Time:	20:41:36	Log-Likelihood:	-728.06
No. Observations:	1380	AIC:	1466.
Df Residuals:	1375	BIC:	1492.
Df Model:	4		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
constant	1.5510	0.073	21.382	0.000	1.409	1.693
hyrsed	0.0669	0.004	17.814	0.000	0.060	0.074
age	0.0144	0.001	12.902	0.000	0.012	0.017
Black	-0.1639	0.045	-3.638	0.000	-0.252	-0.076
Others	0.0307	0.069	0.447	0.655	-0.104	0.165

Omnibus:	96.384	Durbin-Watson:	1.819
Prob(Omnibus):	0.000	Jarque-Bera (JB):	217.837
Skew:	0.424	Prob(JB):	4.98e-48
Kurtosis:	4.752	Cond. No.	291.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS; year == 1980

```

=====
                        OLS Regression Results
=====
Dep. Variable:          ln_hrwage      R-squared:               0.169
Model:                  OLS            Adj. R-squared:         0.167
Method:                 Least Squares   F-statistic:             94.08
Date:                   Mon, 30 Sep 2019 Prob (F-statistic):       6.41e-73
Time:                   20:41:36        Log-Likelihood:          -1148.4
No. Observations:       1856           AIC:                    2307.
Df Residuals:           1851           BIC:                    2334.
Df Model:                4
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
constant	1.6131	0.075	21.590	0.000	1.467	1.760
hyrsed	0.0676	0.004	15.860	0.000	0.059	0.076
age	0.0127	0.001	12.281	0.000	0.011	0.015
Black	-0.1027	0.044	-2.355	0.019	-0.188	-0.017
Others	0.0135	0.071	0.190	0.849	-0.126	0.153

```

=====
Omnibus:                 109.135      Durbin-Watson:           1.958
Prob(Omnibus):           0.000        Jarque-Bera (JB):        266.486
Skew:                    0.337         Prob(JB):                1.36e-58
Kurtosis:                 4.730        Cond. No.                 302.
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

~~~~~
OLS; year == 1990
=====

```

```

                        OLS Regression Results
=====
Dep. Variable:          ln_hrwage      R-squared:               0.217
Model:                  OLS            Adj. R-squared:         0.216
Method:                 Least Squares   F-statistic:             139.3
Date:                   Mon, 30 Sep 2019 Prob (F-statistic):       3.67e-105
Time:                   20:41:36        Log-Likelihood:          -1393.9
No. Observations:       2013           AIC:                    2798.
Df Residuals:           2008           BIC:                    2826.
Df Model:                4
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
constant	1.1186	0.084	13.312	0.000	0.954	1.283

hyrsed	0.0976	0.005	19.991	0.000	0.088	0.107
age	0.0135	0.001	10.785	0.000	0.011	0.016
Black	-0.1720	0.048	-3.601	0.000	-0.266	-0.078
Others	-0.0597	0.089	-0.670	0.503	-0.234	0.115

Omnibus:	111.834	Durbin-Watson:	2.006
Prob(Omnibus):	0.000	Jarque-Bera (JB):	253.156
Skew:	0.342	Prob(JB):	1.07e-55
Kurtosis:	4.597	Cond. No.	349.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS; year == 2000

OLS Regression Results

Dep. Variable:	ln_hrwage	R-squared:	0.207
Model:	OLS	Adj. R-squared:	0.206
Method:	Least Squares	F-statistic:	168.3
Date:	Mon, 30 Sep 2019	Prob (F-statistic):	3.93e-128
Time:	20:41:36	Log-Likelihood:	-2069.0
No. Observations:	2580	AIC:	4148.
Df Residuals:	2575	BIC:	4177.
Df Model:	4		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
constant	1.1617	0.081	14.419	0.000	1.004	1.320
hyrsed	0.1092	0.005	21.168	0.000	0.099	0.119
age	0.0110	0.001	9.585	0.000	0.009	0.013
Black	-0.2460	0.048	-5.145	0.000	-0.340	-0.152
Others	-0.0607	0.060	-1.018	0.309	-0.178	0.056

Omnibus:	305.903	Durbin-Watson:	1.994
Prob(Omnibus):	0.000	Jarque-Bera (JB):	708.951
Skew:	0.696	Prob(JB):	1.13e-154
Kurtosis:	5.158	Cond. No.	339.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[16]: "\nyears = ['data1971.npy', 'data1980.npy', 'data1990.npy', 'data2000.npy',
'data.npy']\nfor t in years:\n    print('OLS for', t)\nprint('=====\n    print(statmod.OLS(endog=t['ln_hrwage'],
exog=t[['constant', 'hyrsed', 'age', 'Black', 'Others']]).fit().summary())\nprint('-----\n-----')\n"
```

Interpretation (Part 4)

To answer this question that how the returns to education change over time in these data, we can note that it is increasing. The results for OLS estimation in the full sample indicate that 1 unit increase in the years of education of males will result in a 7.8 % increase in their wages. Using data for 1971, we can see that this increase in the wage is as much as 6.7% as a result of each year of increased education for the male heads. We can also see that the percentage of increase in the wages is 6.8%, 9.8%, & 11% for 1980, 1990, and 2000, respectively, after one more year of education.

In addition, I estimated the same linear model using the maximum likelihood model and the estimated coefficients for the education are 7.8, 6.97, 6.7, 11.8, 10.9 percent for the full sample, 1971 sample of the data, 1980, 1990 and 2000 samples, respectively, which supports the estimation of the model using OLS both in magnitude and signature. MLE method can be estimated using several optimization methods for the likelihood function. After using Nelder-Mead method and gaining aforementioned estimations, we used bounded L-BFGS-B and SLSQP methods. The former estimates the coefficients of the education as 7.8, 6.7, 6.7, 9.7, and 10.9 % for the full dataset, 1971, 1980, 1990, 2000 datasets respectively which is ageing very close to the estimations of the model using OLS methodology. But the latter estimated far away coefficients for the provided samples of the data with initial values of 0.1 for all coefficients and tolerance level of 1e-15. However, after changing the initial values of the coefficients in the optimization method, I received the same estimations as the ones in OLS and other MLE optimization methods.

All in all, education and wages are positively correlated and increasing the years of education will lead to higher wages.

```
[ ]:
```